Compiler Theory and Practice

Coursework

Juan Scerri

123456A

April 6, 2024

A coursework submitted in fulfilment of study unit CPS2000.

Contents

Listings

Report

1 | Lexer

1.1 | Design & Implementation

The lexer was split into three-main components. A DFSA class, a generic table-driven lexer, and a lexer builder.

The DFSA

The DFSA class is an almost-faithful implementation of the formal concept of a DFSA. Listing 1, outlines the behaviour of the DFSA. Additionally it contains a number of helper functions which facilitate getting the initial state and checking whether a state or a transition category is valid. These helpers specifically, `getInitialState()` is present since after building the DFSA there is no guarantee the initial state used by the user will be the same.

```cpp
class Dfsa {
    public:
     Dfsa(
         size_t noOfStates,
         size_t noOfCategories,
         std::vector<std::vector<int>> const&
             transitionTable,
         int initialState,
         std::unordered_set<int> const& finalStates
    );

    [[nodiscard]] int getInitialState() const;

    [[nodiscard]] bool isValidState(int state) const;
```

```cpp
    [[nodiscard]] bool isValidCategory(int category) const;

    [[nodiscard]] bool isFinalState(int state) const;

    [[nodiscard]] int getTransition(
        int state,
        std::vector<int> const& categories
    ) const;

  private:
    const size_t mNoOfStates;       // Q
    const size_t mNoOfCategories;   // Sigma
    const std::vector<std::vector<int>>
        mTransitionTable;                          // delta
    const int mInitialState;                       // q_0
    const std::unordered_set<int> mFinalStates;   // F
};
```

Listing 1: DFSA Class Declaration (lexer/DFSA.hpp)

The only significant difference is the `getTransition()` functions. In fact, it accepts a vector of transition categories instead of a single category.

This is because a symbol e.g. 'a', '9' etc, might be valid for multiple categories. For instance 'a' is considered to be both a letter and a number in hexadecimal.

The DFSA for accepting the micro-syntax PArL is built as follows.

Let $\mathfrak{U}$ be the set of all possible characters under the system encoding (e.g. UTF-8).

The will use the following categories:

- $L \coloneqq \{\mathrm{A}, \ldots, \mathrm{Z}, \mathrm{a}, \ldots, \mathrm{z}\}$

- $D \coloneqq \{0, \ldots, 9\}$

- $H \coloneqq \{\mathrm{A}, \ldots, \mathrm{F}, \mathrm{a}, \ldots, \mathrm{f}\} \cup D$

- $S \coloneqq \{\alpha \in \mathfrak{U} \colon \alpha \text{ is whitespace}\} \setminus \{\mathrm{LF}\}$

Note: LF refers to line-feed or as it is more commonly known '\n' i.e. new-line.

Together these categories form our alphabet Σ:

2

$$\Sigma := L \cup D \cup S \cup \{\,.\,,\#,\_,(,),[,],\{,\},*,/,+,-,<,>,=,!,\,,\,,:,;,\mathtt{LF}\}$$

Now, the following drawing describe the transitions of the DFSA. For improved readability the DFSA has been split across mulitple drawings. Hence, in each drawing initial state 0 refers to the *same* initial state (a DFSA has one and only one initial state).

Additionally, each final state is annotated with the token type it should produce.

Figure 1: States & transitions for recognising whitespace

Figure 2: States & transitions for recognising identifiers/keywords

Figure 3: States & transitions for recognising builtins

Token::Type::MINUS

Token::Type::ARROW

Figure 4: States & transitions for recognising minus and arrow (->)

Token::Type::INTEGER   Token::Type::FLOAT

Figure 5: States & transitions for recognising integers and floats

Token::Type::COLOR

Figure 6: States & transitions for recognising colours

Figure 7: States & transitions for recognising slashes and comments

Figure 8: States & transitions for assign and is equal to

Figure 9: States & transitions for not equal to

Figure 10: States & transitions for less than and less than or equal to

Token::Type::GREATER

Token::Type::GREATER_EQUAL

Figure 11: States & transitions for greater than and greater than or equal to

Token::Type::LEFT_BRACE   Token::Type::STAR

Token::Type::RIGHT_BRACE   Token::Type::LEFT_PAREN

Token::Type::COLON   Token::Type::RIGHT_PAREN

Token::Type::COMMA   Token::Type::LEFT_BRACK

Token::Type::SEMICOLON   Token::Type::RIGHT_BRACK

Token::Type::PLUS
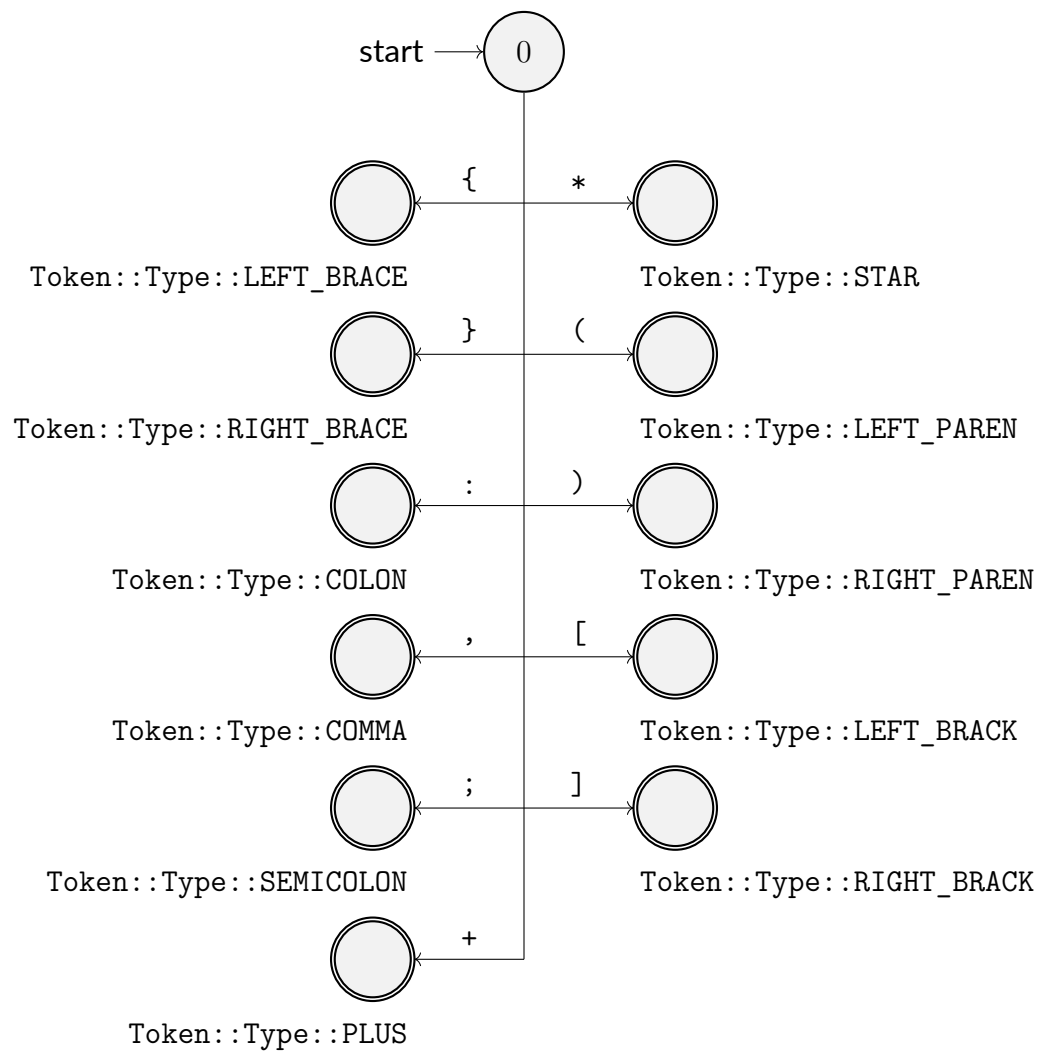
Figure 12: States & transitions for single letter tokens

The Builder & Director

Each sequence of states present is directly represented in code within the `LexerDirector` using methods provided by the `LexerBuilder`.

```
        .addTransition(34, STAR, 36)
        .addComplementaryTransition(36, STAR, 36)
        .addTransition(36, STAR, 37)
        .addComplementaryTransition(37, SLASH, 36)
        .addTransition(37, SLASH, 38)
        .setStateAsFinal(38, Token::Type::COMMENT);

    // builtin
    builder.addTransition(0, UNDERSCORE, 39)
        .addTransition(39, UNDERSCORE, 40)
        .addTransition(40, LETTER, 41)
        .addTransition(41, {LETTER, DIGIT, UNDERSCORE}, 41)
        .setStateAsFinal(41, Token::Type::BUILTIN);
```

Listing 2: Code specification of the comments in the `LexerDirector` (lexer/LexerDirector.cpp)

The `LexerBuilder` keeps track of these transitions using less efficient data structures such as hash maps (`std::unordered_map`) and sets (`std::unordered_set`).

Then the `build()` method processes the user defined transitions and normalises everything into a single transition table for use in a DFSA. Additionally, it also produces two other artefacts. The first is called `categoryIndexToChecker`. It is a hash map from the index of a category to a lambda function which takes a character as input and returns true or false.

The lambdas and the category indices are also registered by the user. See Listing 3 for a registration example. Additionally, the category indices although they are integers for readability they are defined as an enumeration.

```
        .addCategory(
            HEX,
            [](char c) -> bool {
                return ('0' <= c && c <= '9') ||
                       ('A' <= c && c <= 'F') ||
                       ('a' <= c && c <= 'f');
            }
```

7

```
        )
```

Listing 3: Registration of the hexadecimal category checker (lexer/LexerDirector.cpp)

The second artefact produced by the builder is also a hash map from final states to their associated token type.

The transition table is then passed onto the DFSA. And the DFSA, and the two artefacts are passed onto the Lexer class.

```
// create dfsa
Dfsa dfsa(
    noOfStates,
    noOfCategories,
    transitionTable,
    initialStateIndex,
    finalStateIndices
);

// create lexer
Lexer lexer(
    std::move(dfsa),
    std::move(categoryIndexToChecker),
    std::move(finalStateIndexToTokenType)
);
```

Listing 4: Constructions of the Lexer (lexer/LexerBuilder.cpp)

The Actual Lexer

The lexer's core is as was described during the lectures and the core/main method is `simulateDFSA()`.

It also has a number of very important auxiliary methods and behavioural changes. Specifically, the `updateLocationState()`, see Listing 5, is critical for providing adequate error messages both during the current stage and for later stages. This function is called every time a lexeme is consumed allowing the lexer to keep track of where in the file it is, in terms of lines and columns.

```
void Lexer::updateLocationState(std::string const& lexeme) {
```

```
    for (char ch : lexeme) {
        mCursor++;

        if (ch == '\n') {
            mLine++;

            mColumn = 1;
        } else {
            mColumn++;
        }
    }
}
```

Listing 5: The `updateLocationState()` lexer method (lexer/Lexer.cpp)

Additionally, if an invalid / non-accepting state is reached the invalid lexeme is consumed and the user is warned, see Listing 6. After this the lexer, is left in a still operational state. Hence, `nextToken()` can be used again.

This is critical to provide users of the PArL compiler with a list of as many errors as possible, since it would be a bad experience to have to constantly run the PArL compiler to see the next error.

```
    if (state == INVALID_STATE) {
        mHasError = true;

        fmt::println(
            stderr,
            "lexical error at {}:{}:: unexpected "
            "lexeme '{}'",
            mLine,
            mColumn,
            lexeme
        );
    } else {
        try {
            token = createToken(
                lexeme,
                mFinalStateToTokenType.at(state)
            );
        } catch (UndefinedBuiltin& error) {
            mHasError = true;
```

```
        fmt::println(
            stderr,
            "lexical error at {}:{}:: {}",
            mLine,
            mColumn,
            error.what()
        );
    }
}
```

Listing 6: Error handling mechanism in the `nextToken()` lexer method (lexer/Lexer.cpp)

Hooking up the Lexer to the Runner

The `Runner` class is the basic structure which connects all the stages of the compiler together together.

In this case the Runner passes in a reference to the lexer into the parser, this allows the parser to request tokens and they are computed on demand improving overall performance. Additionally, this has the benefit of allowing the parsing of larger and multiple files since, the parser is no longer limited by the amount of usable memory, since it does not need to load the whole file.

However, in this case no such optimisation is present.

```
Runner::Runner(bool dfsaDbg, bool lexerDbg, bool parserDbg)
    : mDfsaDbg(dfsaDbg),
      mLexerDbg(lexerDbg),
      mParserDbg(parserDbg),
      mLexer(LexerDirector::buildLexer()),
      mParser(Parser(mLexer)) {
}
```

Listing 7: The Runner constructor passes `mLexer` into the Parser constructor (runner/Runner.cpp)

2 | The Parser

2.1 | Modified EBNF

Some modifications were applied to the original EBNF. Some of the modifications were either motivated by improved user experience, a more uniform mechanism and others to reduce complexity further down the pipeline.

⟨*Letter*⟩   ::= 'A'-'Z' | 'a'-'z'

⟨*Digit*⟩   ::= '0'-'9'

⟨*Hex*⟩   ::= 'A'-'F' | 'a'-'F' | ⟨*Digit*⟩

⟨*Identifier*⟩   ::= ⟨*Letter*⟩ {'_' | ⟨*Letter*⟩ | ⟨*Digit*⟩}

⟨*BooleanLiteral*⟩   ::= 'true' | 'false'

⟨*IntegerLiteral*⟩   ::= ⟨*Digit*⟩ {⟨*Digit*⟩}

⟨*FloatLiteral*⟩   ::= ⟨*Digit*⟩ {⟨*Digit*⟩} '.' ⟨*Digit*⟩ {⟨*Digit*⟩}

⟨*ColorLiteral*⟩   ::= '#' ⟨*Hex*⟩ ⟨*Hex*⟩ ⟨*Hex*⟩ ⟨*Hex*⟩ ⟨*Hex*⟩ ⟨*Hex*⟩

⟨*ArrayLiteral*⟩   ::= '[' [⟨*Epxr*⟩ {',' ⟨*Epxr*⟩}] ']'

⟨*PadWidth*⟩   ::= '__width'

⟨*PadHeight*⟩   ::= '__height'

⟨*PadRead*⟩   ::= '__read' ⟨*Epxr*⟩ ',' ⟨*Epxr*⟩

⟨*PadRandomInt*⟩   ::= '__random_int' ⟨*Epxr*⟩

⟨*Literal*⟩   ::= ⟨*BooleanLiteral*⟩
               | ⟨*IntegerLiteral*⟩
               | ⟨*FloatLiteral*⟩
               | ⟨*ColorLiteral*⟩
               | ⟨*ArrayLiteral*⟩
               | ⟨*PadWidth*⟩
               | ⟨*PadHeight*⟩
               | ⟨*PadRead*⟩
               | ⟨*PadRandomInt*⟩

⟨*Type*⟩   ::= ('bool' | 'int' | 'float' | 'color') ['[' ⟨*IntegerLiteral*⟩ ']']

⟨*SubEpxr*⟩	::=	'(' ⟨*Epxr*⟩ ')'
⟨*Variable*⟩	::=	⟨*Identifier*⟩
⟨*ArrayAccess*⟩	::=	⟨*Identifier*⟩ '[' ⟨*Epxr*⟩ ']'
⟨*FunctionCall*⟩	::=	⟨*Identifier*⟩ '(' [⟨*Epxr*⟩ {',' ⟨*Epxr*⟩}] ')'
⟨*Epxr*⟩	::=	⟨*LogicOr*⟩ ['as' ⟨*Type*⟩]
⟨*LogicOr*⟩	::=	⟨*LogicAnd*⟩ {'or' ⟨*LogicAnd*⟩}
⟨*LogicAnd*⟩	::=	⟨*Equality*⟩ {'and' ⟨*Equality*⟩}
⟨*Equality*⟩	::=	⟨*Comparison*⟩ {('==' | '!=') ⟨*Comparison*⟩}
⟨*Comparison*⟩	::=	⟨*Term*⟩ {('<' | '<=' | '>' | '>=') ⟨*Term*⟩}
⟨*Term*⟩	::=	⟨*Factor*⟩ {('+' | '−') ⟨*Factor*⟩}
⟨*Factor*⟩	::=	⟨*Unary*⟩ {('*' | '/') ⟨*Unary*⟩}
⟨*Unary*⟩	::=	('−' | 'not') ⟨*Unary*⟩ | ⟨*Primary*⟩
⟨*RefExpr*⟩	::=	⟨*Variable*⟩
	|	⟨*ArrayAccess*⟩
	|	⟨*FunctionCall*⟩
⟨*Primary*⟩	::=	⟨*Literal*⟩
	|	⟨*SubExpr*⟩
	|	⟨*RefExpr*⟩
⟨*Program*⟩	::=	{⟨*Stmt*⟩}
⟨*Stmt*⟩	::=	⟨*Block*⟩
	|	⟨*VaribaleDecl*⟩ ';'
	|	⟨*FunctionDecl*⟩
	|	⟨*Assignment*⟩ ';'
	|	⟨*PrintStmt*⟩ ';'
	|	⟨*DelayStmt*⟩ ';'
	|	⟨*WriteBoxStmt*⟩ ';'
	|	⟨*WriteStmt*⟩ ';'
	|	⟨*ClearStmt*⟩ ';'
	|	⟨*IfStmt*⟩

	|	⟨*ForStmt*⟩
	|	⟨*WhileStmt*⟩
	|	⟨*ReturnStmt*⟩ ';'

⟨*Block*⟩ ::= '{' {⟨*Stmt*⟩} '}'

⟨*VariableDecl*⟩ ::= 'let' ⟨*Identifier*⟩ ':' ⟨*Type*⟩ '=' ⟨*Epxr*⟩

⟨*FormalParam*⟩ ::= ⟨*Identifier*⟩ ':' ⟨*Type*⟩

⟨*FunctionDecl*⟩ ::= 'fun' ⟨*Identifier*⟩ '(' [⟨*ForamlParam*⟩ {',' ⟨*FormalParam*⟩}] ')'
'->' ⟨*Type*⟩ ⟨*Block*⟩

⟨*Assignment*⟩ ::= ⟨*Identifier*⟩ ['[' ⟨*Epxr*⟩ ']'] '=' ⟨*Epxr*⟩

⟨*PrintStmt*⟩ ::= '__print' ⟨*Epxr*⟩

⟨*DelayStmt*⟩ ::= '__delay' ⟨*Epxr*⟩

⟨*WriteBoxStmt*⟩ ::= '__write_box' ⟨*Epxr*⟩',' ⟨*Epxr*⟩','⟨*Epxr*⟩',' ⟨*Epxr*⟩','⟨*Epxr*⟩

⟨*WriteStmt*⟩ ::= '__write' ⟨*Epxr*⟩',' ⟨*Epxr*⟩','⟨*Epxr*⟩

⟨*ClearStmt*⟩ ::= '__clear' ⟨*Epxr*⟩

⟨*IfStmt*⟩ ::= 'if' '(' ⟨*Expr*⟩ ')' ⟨*Block*⟩ ['else' ⟨*Block*⟩]

⟨*ForStmt*⟩ ::= 'for' '(' [⟨*VariableDecl*⟩] ';' ⟨*Expr*⟩ ';' [⟨*Assignment*⟩] ')' ⟨*Block*⟩

⟨*WhileStmt*⟩ ::= 'while' '(' ⟨*Expr*⟩ ')' ⟨*Block*⟩

⟨*ReturnStmt*⟩ ::= 'return' ⟨*Expr*⟩

Improved Precedence

So, the minor changes which improve programmer usability are the additions of a number of other expression stages, such as ⟨LogicOr⟩, ⟨LogicAnd⟩, etc. The main reason for the addition of such rules is to further enforce a more natural operation precedence. For example a programmer often expects that comparison operators such as < and > bind tighter than and or or, hence the compiler needs to make sure that comparison operators are executed before logical operators, and this can be enforced by the grammar itself hence the changes.

Better Arrays

The way arrays were being implemented in the original grammar was very restrictive. Instead an approach for treating arrays as their own type and literal was taken up.

The ⟨Type⟩ and ⟨Literal⟩ productions were augmented to improve array support. This helped simplify the ⟨Identifier⟩ (in the original EBNF) , ⟨VariableDecl⟩ and ⟨FormalParam⟩ productions. Additionally, this opens up further support for more complicated types later on.

For example, the ⟨Type⟩ can be further augmented to support more types.

⟨StructField⟩ ::= ⟨Identifier⟩ ':' ⟨Type⟩

⟨Struct⟩ ::= 'struct' ⟨Identifier⟩ '{' {⟨StructField⟩ ';'} '}'

⟨TypeDecl⟩ ::= ⟨Struct⟩

⟨Base⟩ ::= ('bool' | 'int' | 'float' | 'color')

⟨Array⟩ ::= ⟨Type⟩ '[' ⟨IntegerLiteral⟩ ']'

⟨Pointer⟩ ::= ⟨Type⟩ '*'

⟨Type⟩ ::= ⟨Identifier⟩
| ⟨Base⟩
| ⟨Array⟩
| ⟨Pointer⟩

Note that we are indeed repeating ⟨ForamlParam⟩ but this is not really a problem. When it comes to specification repetition which improves clarity is "good" repetition.

Additionally, some of the ground work for these improvements has already been laid out in the internal type system see Listing 33 and Listing 9.

```
struct Primitive;

enum class Base {
    BOOL,
    COLOR,
    FLOAT,
```

```
    INT,
};

struct Array {
    size_t size;
    box<Primitive> type;
};
```

Listing 8:   Mechanism for internally storing types within the compiler (parl/Core.hpp)

```
struct Primitive {
    template <typename T>
    [[nodiscard]] bool is() const {
        return std::holds_alternative<T>(data);
    }

    template <typename T>
    [[nodiscard]] const T &as() const {
        return std::get<T>(data);
    }

    ...

    bool operator!=(Primitive const &other) {
        return !operator==(other);
    }

    std::variant<std::monostate, Base, Array> data{};
};
```

Listing 9: The Primitive structure (parl/Core.hpp)

The box type is a special type of pointer object which has value semantics that is it behaves as though it were the object it contains.

This is critical because self-referential types like the ⟨Array⟩ as described above are not easily representable within C++ since, something like Listing 10, is not allowed.

```
struct SelfRef;
```

```cpp
struct Container {
    Other other;
    SelfRef ref;
};

struct Primitive {
    std::variant<std::monostate, Container> data{};
};
```

Listing 10: An size unbounded type in C++ (parl/Core.hpp)

This is because the C++ compiler is incapable of determining the size of said type at compile-time. Hence, a pointer for said type is required. And the pointer wrapper `box<>` allows it to be copied as tough it were a value.

The `box<>` type is attributed to Jonathan Müller were he discuss the exact issue described above on his blog, `foonathan`.

These changes to type also require additional, changes to how variables are referenced in expressions hence why ⟨RefExpr⟩ was added. It also ensure that in the future more complicated referencing such as `object.something` (`struct` member referencing) is possible.

Finally, the ⟨ArrayLiteral⟩ has been improved to support expressions instead of just only literals.

Removing Eye-Candy

Adding this system however, significantly increases the complexity of semantic analysis.

Because of this the slight syntax sugar of being able to have the following two conveniences `let a: int[] = [1,1,1];` and `let a: int[3] = [1];` has been dropped.

This is because such syntax not only complicates the grammar but it also significantly increases the complexity of semantic analysis.

The best way to handle this is to have a de-sugaring & type inference sub-phase before type checking, within the semantic analysis phase.

2.2 | Parsing & The Abstract Syntax Tree

The AST is the data structure which is produced by the parser. The nodes of
the AST are in fact almost a one-to-one representation of the productions in the
grammar, for example see Listing 11.

```cpp
    void accept(Visitor*) override;

    std::unique_ptr<Expr> subExpr;
};

struct Binary : public Expr {
    explicit Binary(std::unique_ptr<Expr>, Operation, std::::
    unique_ptr<Expr>);
```

Listing 11: The `FunctionCall` AST node class (parl/AST.hpp)

The only significant difference in these nodes in the `Position` field. This get pop-
ulated by the parser using the location of a token in the original source file. This
is again critical for adequate error messaging later in the semantic analysis phase.

```cpp
    void accept(Visitor*) override;

    const Operation op;
    std::unique_ptr<Expr> expr;
};

struct Stmt : public Node {};

struct Assignment : public Stmt {
```

Listing 12: The `Binary` AST node class (parl/AST.hpp)

```cpp
    void accept(Visitor*) override;

    const std::string identifier;
    std::unique_ptr<Expr> index;
    std::unique_ptr<Expr> expr;
```

17

```
};

struct VariableDecl : public Stmt {
```

Listing 13: The `Unary` AST node class (parl/AST.hpp)

Apart from these the only other real difference is the use of a `Binary` and `Unary` node see Listing 12 and Listing 13, instead of a node for each type of the grammar rules discussed above in 2.1. This is because the reason for said rules was precedence and precedence within the AST is actually described by the structure of the tree itself not the node types.

Additionally, a number of the nodes override the `accept()` method specified by the pure virtual class `Node`. This is the basis for the Visitor pattern which apart form the parser is the backbone of the remaining phases.

2.3 | The Actual Parser

The parser is split into four main sections.

- AST Generation

- Token Buffering

- Token Matching

- Error Handling/Recovery

Token Buffering

The parser requires access to the tokens for proper functioning. Additionally, sometimes the parser requires more than one token as the case deciding on whether an identifier is just a variable or a function call. This is quite easy to implement if the parser has available to it at initialisation all the tokens.

However, as described in 1.1. This is not the case the parser requests token on demand from lexer which it has a reference. This of course means that the machinery for handling tokens is a bit more complicated due to requiring lookahead.

To solve this issue a window-based approach was adopted. The parser has a moving buffer/window called `mTokenBuffer` and whose size is specified at compile-time using a C-style macro `#define LOOKAHEAD (2)`. The core methods for this aspect of the parser are `moveWindow()` and `nextToken()`, see Listing 14.

```
    Token token = previous();

    consume(
        Token::Type::LEFT_PAREN,
        "expected '(' after identifier"
    );

    std::vector<std::unique_ptr<core::Expr>> params{};

    if (!peekMatch({Token::Type::RIGHT_PAREN})) {
        do {
            params.emplace_back(expr());
        } while (match({Token::Type::COMMA}));
    }

    consume(
        Token::Type::RIGHT_PAREN,
        "expected ')' after parameters"
    );
```

Listing 14: The `moveWindow()` and `nextToken()` Parser methods
(parser/Parser.cpp)

Additionally, note that within `nextToken()` the whitespace and comments are
being explicitly ignored.

This might lead to further minor improvement. Specifically, comments can be
integrated into the AST. This basically allows printing visitors or, formatting vis-
itors properly format the code whilst still preserving any comments created by
programmers.

Token Matching

The previous methods all facilitate the more important token matching methods
which are:

- `peek()`,

- `advance()`,

- `previous()`,

- isAtEnd(),

- peekMatch(),

- match(),

- and, consume().

Arguably, the most improtant of these methods is `consume()`, it takes in a token type and a error message if the specified token type is not matched, see Listing 15.

```
template <typename... T>
void consume(
    Token::Type type,
    fmt::format_string<T...> fmt,
    T&&... args
) {
    if (check(type)) {
        advance();
    } else {
        error(fmt, args...);
    }
}
```

Listing 15: The `consume()` Parser methods (parser/Parser.hpp)

The main reason for the templating of such a method is to provide an easy interface for formattable strings using fmtlib. The main feature this library provides is the ability to specify placeholders in the string itself using {}. Usage of this functionality is demonstrated in the AST generator methods, see Listing 16.

```
        Token::Type::RIGHT_PAREN,
        "expected ')' after expression"
    );

    std::unique_ptr<core::Block> block_ = block();
```

Listing 16: The Usage of the `consume()` method in the `formalParam()` generator method (parser/Parser.cpp)

The `peekMatch()` method is a simple method which returns true if at least one of the provided token types match. The main usage for this is cases where more

than different token is valid option for example such is the case for ⟨Comparison⟩ production. `match()` is a simple extension of `peekMatch()` which consumes the token if it matches.

The methods `advance()`, `previous()` and `isAtEnd()` are quite self explanatory. `advance()` moves the buffer window one step forward, `previous()` returns the last consumed token, and `isAtEnd()` checks whether or not the parser has reached an End of File token.

`peek()` is also very simple it allows the parser to see the token without consuming it. However, since the parser might need to lookahead peek() supports an offset. The only issue is that all calls to `peek()` have to be checked to ensure that valid access, see Listing 17. This is done using `abort_if()` function which prints an error message an aborts the program. This is very similar to a `static_assert` however, it is has to be performed at runtime, see Listing 18.

```cpp
}

void Parser::initWindow() {
    for (int i = 0; i < LOOKAHEAD; i++) {
        mTokenBuffer[i] = nextToken();
    }

    mPreviousToken = mTokenBuffer[0];
}
```

Listing 17: The `peek()` Parser method (parser/Parser.cpp)

```cpp
#ifdef NDEBUG
template <typename... T>
inline void abort(fmt::format_string<T...>, T &&...) {
}

template <typename... T>
inline void
abort_if(bool, fmt::format_string<T...>, T &&...) {
}
#else
template <typename... T>
inline void
abort(fmt::format_string<T...> fmt, T &&...args) {
    fmt::println(stderr, fmt, args...);
```

```cpp
    std::abort();
}

template <typename... T>
inline void abort_if(
    bool cond,
    fmt::format_string<T...> fmt,
    T &&...args
) {
    if (cond) {
        abort(fmt, args...);
    }
}
#endif
```

Listing 18: The abort functionality present in the codebase (parl/Core.hpp)

Note that since the usage of `abort()` and `abort_if()` is internal to the code, it only makes sense for these functions to be enabled during debug builds only. Hence, the implementation are enclosed between an `#ifdef`, `#else`, `#endif` macro. When `NDEBUG` which stand for no-debug is defined the function bodies are hollowed out allowing the C++ compiler to optimise them out.

Error Handling/Synchronization

Error handling and synchronization is a critical part of the parser. With regards to developer productivity, having meaningful errors is extremely valuable. But apart from that being able see all the errors in a file is also crucial. This means that a developer will waste less time re-running the compiler to find all the errors present in the source code. This processes is referred to as Synchronization by the author of Crafting Interpreters, Robert Nystrom. The method for synchronization in the PArL compiler was inspired by the above credited author and his usage of Exceptions as an unrolling primitive although unorthodox is very simple to implement and has in fact been used in the parser and also later stages of the compiler with success.

```cpp
class SyncParser : public std::exception {};

...
```

```cpp
template <typename... T>
void error(fmt::format_string<T...> fmt, T&&... args) {
    mHasError = true;

    Token violatingToken = peek();

    fmt::println(
        stderr,
        "parsing error at {}:{}:: {}",
        violatingToken.getPosition().row(),
        violatingToken.getPosition().col(),
        fmt::format(fmt, args...)
    );

    throw SyncParser{};
}
```

Listing 19: The `ParseSync` exception and the `error()` method which kickstarts the synchronization process (parser/Parser.hpp)

However, there is of course a major downside to this where synchronization happens will affect other error messages which are reported down stream. This of course has the possibility of producing false positives. However, in this case error handling is only best-effort and therefore the possibility of false positives accepted. The basic process of synchronizing is consuming as many tokens as possible until the parser reaches a token which it believes to be a good restarting point, see Listing 20.

```cpp
}

bool Parser::check(Token::Type type) {
    if (isAtEnd()) {
        return false;
    }

    return peek().getType() == type;
}

bool Parser::peekMatch(
    std::initializer_list<Token::Type> const &types
) {
    return std::any_of(
```

```cpp
        types.begin(),
        types.end(),
        [&](auto type) {
            return check(type);
        }
    );
}

bool Parser::match(
    std::initializer_list<Token::Type> const &types
) {
    if (peekMatch(types)) {
        advance();

        return true;
    }

    return false;
}

// NOTE: synchronization is all best effort

void Parser::synchronize() {
    while (!isAtEnd()) {
        Token peekToken = peek();

        switch (peekToken.getType()) {
            case Token::Type::SEMICOLON:
                advance();

                return;
            case Token::Type::FOR:
                /* fall through */
```

Listing 20: The `synchronize()` method in the Parser class (parser/Parser.cpp)

Finally, the most natural choice for capturing `ParserSync` is in AST generator methods responsible for generating statement nodes, those being `program()` and `block()`.

AST Generator Methods

Finally, the bulk of the parser is actually the generator methods which actually build the AST. These methods are not that complicated and they follow the specified grammar faithfully.

See, Listing 21 for an example of such a method and see Listing 22 for a full list of the methods.

```
    consume(
        Token::Type::COMMA,
        "expected ',' after expression"
    );

    std::unique_ptr<core::Expr> xOffset = expr();

    consume(
        Token::Type::COMMA,
        "expected ',' after expression"
    );

    std::unique_ptr<core::Expr> yOffset = expr();

    consume(
        Token::Type::COMMA,
        "expected ',' after expression"
    );

    std::unique_ptr<core::Expr> color = expr();

    return make_with_pos<core::WriteBoxStmt>(
        token.getPosition(),
        std::move(x),
        std::move(y),
        std::move(xOffset),
        std::move(yOffset),
        std::move(color)
    );
}

std::unique_ptr<core::IfStmt> Parser::ifStmt() {
    consume(
        Token::Type::IF,
```

```
      "expected 'if' at start of if statement"
```

Listing 21: The `ifStmt()` node generator method in the Parser class (parser/Parser.cpp)

```cpp
std::unique_ptr<core::Type> type();

std::unique_ptr<core::Program> program();
std::unique_ptr<core::Stmt> statement();
std::unique_ptr<core::Block> block();
std::unique_ptr<core::VariableDecl> variableDecl();
std::unique_ptr<core::Assignment> assignment();
std::unique_ptr<core::PrintStmt> printStatement();
std::unique_ptr<core::DelayStmt> delayStatement();
std::unique_ptr<core::WriteBoxStmt> writeBoxStatement();
std::unique_ptr<core::WriteStmt> writeStatement();
std::unique_ptr<core::ClearStmt> clearStatement();
std::unique_ptr<core::IfStmt> ifStmt();
std::unique_ptr<core::ForStmt> forStmt();
std::unique_ptr<core::WhileStmt> whileStmt();
std::unique_ptr<core::ReturnStmt> returnStmt();
std::unique_ptr<core::FunctionDecl> functionDecl();
std::unique_ptr<core::FormalParam> formalParam();

std::unique_ptr<core::PadWidth> padWidth();
std::unique_ptr<core::PadHeight> padHeight();
std::unique_ptr<core::PadRead> padRead();
std::unique_ptr<core::PadRandomInt> padRandomInt();
std::unique_ptr<core::BooleanLiteral> booleanLiteral();
std::unique_ptr<core::ColorLiteral> colorLiteral();
std::unique_ptr<core::FloatLiteral> floatLiteral();
std::unique_ptr<core::IntegerLiteral> integerLiteral();
std::unique_ptr<core::ArrayLiteral> arrayLiteral();
std::unique_ptr<core::SubExpr> subExpr();
std::unique_ptr<core::Variable> variable();
std::unique_ptr<core::ArrayAccess> arrayAccess();
std::unique_ptr<core::FunctionCall> functionCall();

std::unique_ptr<core::Expr> expr();
std::unique_ptr<core::Expr> logicOr();
std::unique_ptr<core::Expr> logicAnd();
std::unique_ptr<core::Expr> equality();
```

```
std::unique_ptr<core::Expr> comparison();
std::unique_ptr<core::Expr> term();
std::unique_ptr<core::Expr> factor();
std::unique_ptr<core::Expr> unary();
std::unique_ptr<core::Expr> primary();
```

Listing 22: The main body of methods in the Parser class (parser/Parser.hpp)

2.4 | Pretty? Printing

Using Parser in the Runner

The parser is used in the runner and assuming that the parser does not encounter any errors and the `mParserDbg` flag is set it can be used to print the AST using a `PrinterVisitor`.

```
mParser.parse(source);

if (mLexer.hasError() || mParser.hasError()) {
    return;
}

std::unique_ptr<core::Program> ast = mParser.getAst();

if (mParserDbg) {
    debugParsing(ast.get());
}
```

Listing 23: The parse segment of the `run()` method in the Runner class (runner/Runner.cpp)

Calling the produced `PArL` binary with the `-p` flag will set the `mParserDbg`, see Figure 13 and Figure 14.

```
> cat testing/test9.parl
fun AverageOfTwo(x: int, y: int) -> float {
    let t0: int = x + y;
    if (t0 > 10) {
        return 10;
    }
    let t1: float = t0 / 2 as float;
    return t1;
}
```

Figure 13: `cat` of the test9.parl

```
> ./run.sh -p testing/test9.parl
Parser Debug Print
Program =>
  Func Decl AverageOfTwo =>
    Formal Param x =>
      int
    Formal Param y =>
      int
    float
    {
      let t0 :
        int
        Binary Operation + =>
          Variable x
          Variable y
      If =>
        Binary Operation > =>
          Variable t0
          int 10
        {
          Return =>
            int 10
        }
      let t1 :
        float
        Binary Operation / =>
          Variable t0
          int 2
        as
          float
      Return =>
        Variable t1
    }
semantic error at 4:9:: incorrect return type in function AverageOfTwo
```

Figure 14: The AST generated by the syntactically correct program in test9.parl

```cpp
void Runner::debugParsing(core::Program* program) {
    fmt::println("Parser Debug Print");

    PrinterVisitor printer;

    program->accept(&printer);
}
```

Listing 24: The `debugParsing()` method in the Runner class (runner/Runner.cpp)

The `PrinterVisitor` used in Listing 24 is a specialization of the pure virtual `Visitor` class in parl/Visitor.hpp.

```
...

struct ClearStmt;
struct Block;
struct FormalParam;
struct FunctionDecl;
struct IfStmt;
struct ForStmt;
struct WhileStmt;
struct ReturnStmt;
struct Program;

class Visitor {
  public:
    virtual void visit(Type*) = 0;
    virtual void visit(Expr*) = 0;
    virtual void visit(PadWidth*) = 0;
    virtual void visit(PadHeight*) = 0;
    virtual void visit(PadRead*) = 0;
    virtual void visit(PadRandomInt*) = 0;
    virtual void visit(BooleanLiteral*) = 0;
    virtual void visit(IntegerLiteral*) = 0;
    virtual void visit(FloatLiteral*) = 0;

...
```

Listing 25: A segment of the pure virtual `Visitor` class (parl/Visitor.hpp)

Due to the way C++ handles symbols, the classes which the visitor can visit must be forward declared manually, see Listing 25. If no such forward declaration is made, the compiler will complain about the `Visitor` and the `AST` classes being cyclically dependent.

Additionally, any inheriting visitor such as the `PrinterVisitor` can hold state. In fact, this is where the true power of visitors arises. Being able to hold state

means that complex computations can be carried out on the AST. For example the `PrinterVisitor` although simple makes use of a single variable `mTabCount`, see Listing 26, which it uses to affect how much indentation should be used in printing, allowing us to visualise the AST.

```cpp
void PrinterVisitor::visit(core::PadRead *expr) {
    print_with_tabs("__read =>");
    mTabCount++;
    expr->x->accept(this);
    expr->y->accept(this);
    mTabCount--;

    expr->core::Expr::accept(this);
}
```

Listing 26: The `visit(core::PadRead*)` method in the `PrinterVisitor` (parser/PrinterVisitor.cpp)

3 │ Semantic Analysis

3.1 │ Phases & Design

As described in 2.1 the Semantic Analysis phase should be split into a number of sub-phases specifically: symbol resolution, de-sugaring/type inference and type checking.

However, these steps can be combined together into a single step phase. There a benefits and downsides to both approaches. The main benefit of using the first approach is a reduction in algorithmic complexity. The logic can be separated into different phases with ease and information from one phase can be propagated to another. The downside of this approach is that it actually adds more code for maintenance, since each individual sub-phase will probably need to be implemented as its own visitor. The other down side is error management. If an error occurs in a particular phase since said phase is completely disjoint from the phases after it a mechanism for propagation needs to be devised which again further increases complexity. Of course by the very nature of this argument the monolithic approach does not suffer for the issues that the sub-phases approach has. However, it significantly increases complexity since all sub-phases are being done in a single phase. The other more glaring issue is the fact that it is much more difficult to resolve symbols before-hand.

You would want to do so to allow for the location of function declarations in code to not effect resolution, that is, a function can be called before it is referenced.

Unfortunately, due to the fact that compiler development was an organic process and not too much time was spend on deliberation the semantic analysis phase became monolithic, and hence it suffers from the issues described above. Of course, this means location agnostic function declaration are currently *not* supported by the compiler.

3.2 | The Environment Tree

Some terminology is required to properly describe the number of structures which will be used in this section. The following terms are critical and need to be differentiated properly:

- `SymbolTable`

- `SymbolTableStack`

- `Environment`

- `EnvStack`

- `RefStack`

During semantic analysis there is a need for specific data structures which facilitate scoping rules, type checking, etc.

The most basic data structure which achieves this, is a single `SymbolTable`. A symbol table in its simplest form is a wrapper around a hash map, whose keys are identifiers and values are `Symbols` or any structure which is capable of storing variable and function signatures.

This approach is quite limiting since it restricts developers and user of the language to a single global scope. Therefore, this is structure is not a sufficient solution for most toy-languages let alone production ready languages.

A better solution is using a stack of symbol tables. This allows symbol tables to shadow each other allowing for the reuse of identifiers. Additionally, this further opens up the possibility of implementing modules at the language-level. The likelihood that names will be reused across different modules is quite high and being able to scope modules so they do not interfere with global scope and each other is necessary for any sufficiently large project.

Figure 15: A PArL function with annotated scopes

The basic premise of using a symbol table stack is described in Algorithm 1 and Figure 16. A new symbol table, also referred to as a scope, is pushed onto a stack only for specific types of AST nodes. The node is then processed, were "processing" often times means considering all the sub-children of the node, and when processing terminates the scope is popped. In the context of a purely stack based implementation this implies that, scopes are only temporary that is they are lost after being popped.

Algorithm 1: Basic description of `SymbolTableStack` usage

Data: N the AST node, S the symbol table stack
begin
    if N *opens a scope* **then**
        PushScope(S);
    end
    ProcessNode(N);
    if N *opens a scope* **then**
        PopScope(S);
    end
end
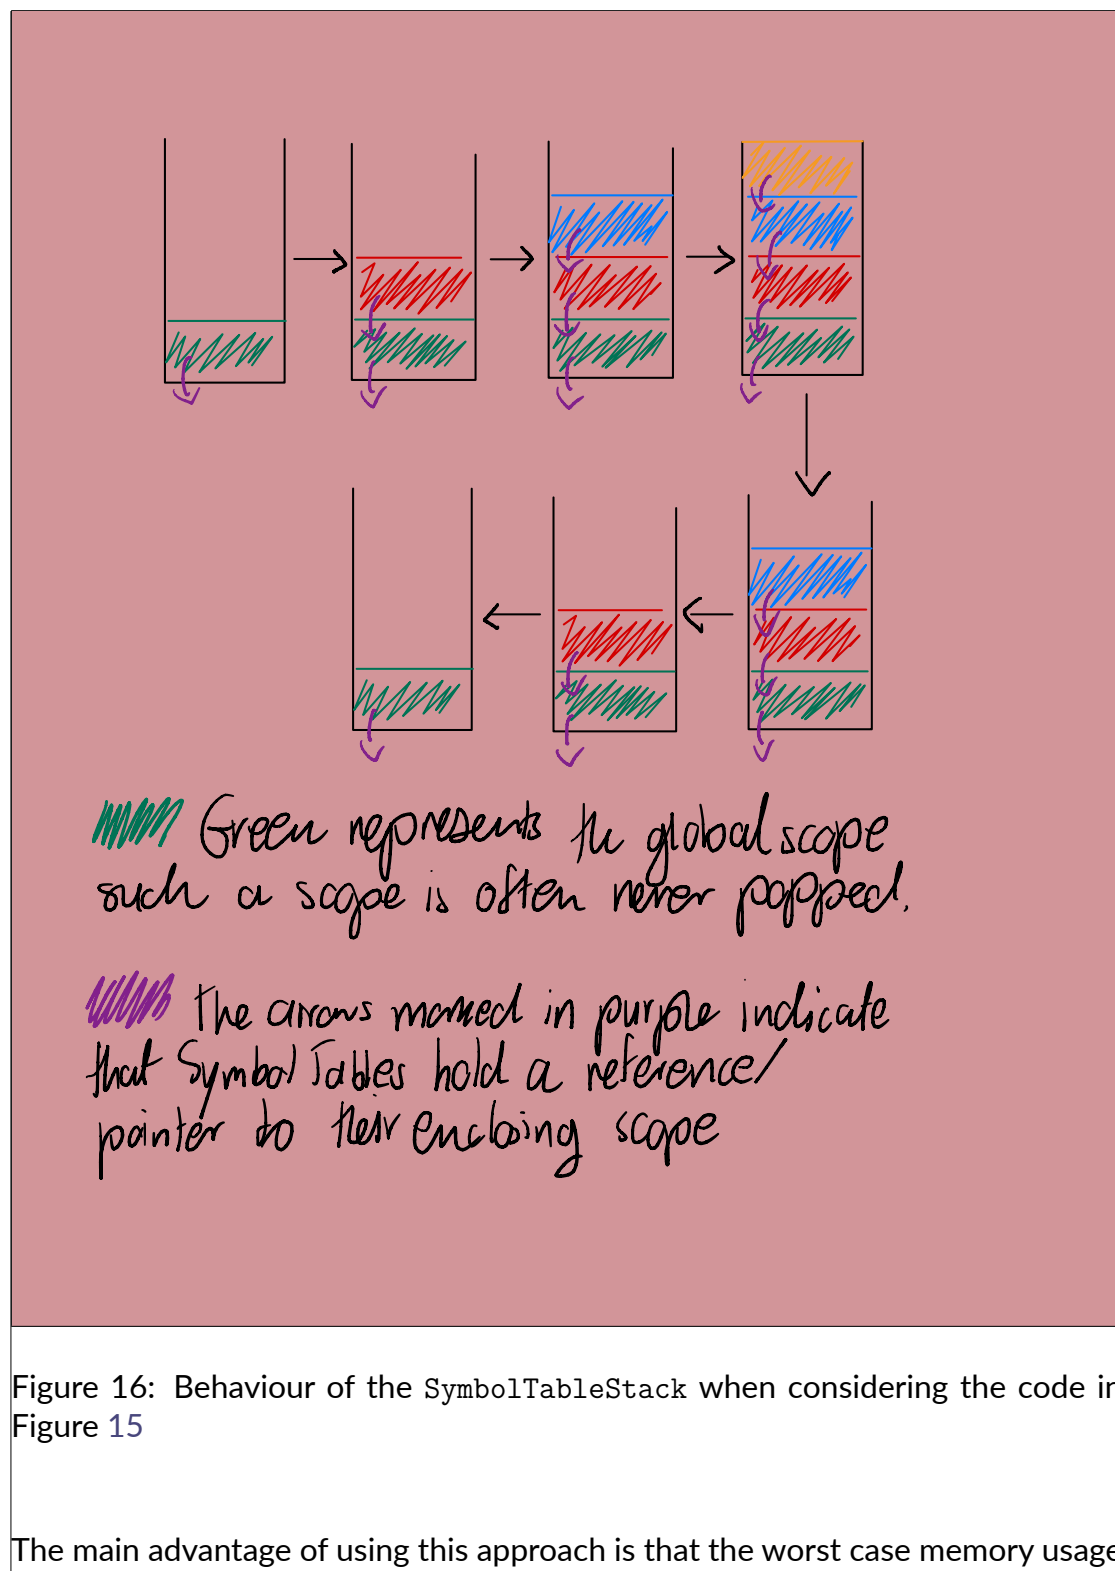
is $O(DI)$ where D is the length of a longest chain of open scopes and I is the size of largest number of variable declarations within a scope.

However, due to the temporary nature of this approach it does not lend itself well to a multi-phase approach. This is because at each phase `SymbolTables` have to be regenerated increase the complexity of each individual phase.

Instead a different approach shall be used. This approach uses an Environment Tree and it has also been inspired by Crafting Interpreters.

The notion of an `Environment` is essentially, the same as a `SymbolTable`. The only significant change is the addition of vector of unique pointers to other environments, see Listing 27.

```cpp
class Environment {
   public:
    enum class Type {
        GLOBAL,
        IF,
        ELSE,
        FOR,
        WHILE,
        FUNCTION,
        BLOCK
    };

    void addSymbol(
        std::string const& identifier,
        Symbol const& Symbol
    );
    [[nodiscard]] std::optional<Symbol> findSymbol(
        std::string const& identifier
    ) const;
    Symbol& getSymbolAsRef(std::string const& identifier);

    [[nodiscard]] Environment* getEnclosing() const;
    void setEnclosing(Environment* enclosing);

    [[nodiscard]] Type getType() const;
    void setType(Type type);

    [[nodiscard]] std::optional<std::string> getName(
    ) const;
```

```cpp
    void setName(const std::string& name);

    std::vector<std::unique_ptr<Environment>>& children();

    [[nodiscard]] bool isGlobal() const;

    [[nodiscard]] size_t getIdx() const;
    void incIdx();
    void incIdx(size_t inc);

    [[nodiscard]] size_t getSize() const;
    void setSize(size_t size);

  private:
    std::unordered_map<std::string, Symbol> mMap{};
    Type mType{Type::GLOBAL};
    std::optional<std::string> mName{};
    Environment* mEnclosing{nullptr};
    std::vector<std::unique_ptr<Environment>> mChildren{};
    size_t mSize{0};
    size_t mIdx{0};
};
```

Listing 27: The `Environment` class with `mChildren` highlighted (backend/Environment.hpp)

The additional fields present in the `Environment` class are there to cater for the needs of each of the phases.

Specifically, `mType` and `mName` are used during semantic analysis and `mSize` and `mIdx` are used during code generation.

Additionally, it is preferable if the interface with which the phases interact with the `Environment` tree is identical to that of a `SymbolTableStack`, that is the same `push()` and `pop()` methods are used.

Now to facilitate this another two classes need to be defined `EnvStack` and `RefStack`. The main difference between `EnvStack` and `RefStack` is that the purpose of `EnvStack` is to construct the `Environment` tree whilst `RefStack` traverses the environment tree. Due to this `RefStack` is only required during the symbol resolution. However, in current implementation since symbol resolution is combined with type checking, the Environment tree is generated at the same time as it is being type checked.

Additionally, in both an `EnvStack` and a `RefStack`, the creation and traversal of the Environment tree are managed via the `push()` and `pop()` methods.

In an `EnvStack` `push()` works by creating a new environment, setting the enclosing environment to the current environment and changing the current environment to to the new environment, see Listing **??**. `pop()` makes use of the `mEnclosing` field in the current environment and just sets the current environment to the enclosing environment essentially moving up an environment, see Listing **??**.

```
EnvStack& EnvStack::pushEnv() {
    auto& ref = mCurrent->children().emplace_back(
        std::make_unique<Environment>()
    );

    ref->setEnclosing(mCurrent);

    mCurrent = ref.get();

    return *this;
}
```

Listing 28: The `pushEnv()` method in the `EnvStack` class (analysis/EnvStack.cpp)

```
EnvStack& EnvStack::popEnv() {
    mCurrent = mCurrent->getEnclosing();

    return *this;
}
```

Listing 29: The `popEnv()` method in the `EnvStack` class (analysis/EnvStack.cpp)

The implementations in `RefStack` are however a bit more involved. This is because the program has to be able to perform a step-wise left-first depth-first traversal. The best way to do this is to make use of a stack and the algorithms in Listing 30 and Listing 31.

```
RefStack& RefStack::pushEnv() {
    size_t index = mStack.top();

    mStack.top()++;
```

```
    mStack.push(0);

    mCurrent = mCurrent->children()[index].get();

    return *this;
}
```

Listing 30: The `pushEnv()` method in the `RefStack` class (ir_gen/RefStack.cpp)

```
RefStack& RefStack::popEnv() {
    mStack.pop();

    mCurrent = mCurrent->getEnclosing();

    return *this;
}
```

Listing 31: The `popEnv()` method in the `RefStack` class (ir_gen/RefStack.cpp)

See Figure 17, for the initial segments of a traversal. The important thing to note here is that the correctness of the traversal entirely depends on the usage of `push()` and `pop()`. If used incorrectly the traversal might be incorrect or the program might crash. Of course, this is only an implementation detail that is if used correctly by the compiler developer no such issue should occur.

Figure 17: The initial steps of a traversal performed by `RefStack`

In the current implementation of `PArL` three types of nodes control scopes, these are: `Blocks`, `ForStmts` and `FunctionDecls`. There is also a slight particularity in

that, `ForStmts` and `FunctionDecls` as per the grammar actually, casue two scopes to open. This really depends on whether shadowing of, generally the iterator in a for-loop and the function parameters in the language should be allowed or not. For simplicities sake in PArL a second scope is being opened. However there is no particular consensus, for example in C++ shadowing of function parameters or declarations in for-loops is not allowed since they create a single scope, but Rust allows shadowing for the same variable in the same scope without any problems.

3.3 | Semantic Analysis

Having the necessary data structures in place the discussion will now revolve around Semantic Analysis i.e. the process of making sure the meaning of the program is correct.

There are four main areas of interest which shall be discussed: the symbol type, symbol registration/search, type checking and a unique subset of type checking, return value type verification.

3.4 | The Symbol Type

Within PArL functions are special they are not treated as values i.e. they they cannot be assigned to variables, they cannot be moved around etc. This means that the `Symbol` type is actually larger than the `Primitve` type referenced in 2.1. It has to support both variables and functions, hence the `Symbol` is truly the union of two different types, `VariableSymbol` and `FunctionSymbol`.

```cpp
struct VariableSymbol {
    size_t idx{0};
    core::Primitive type;

    ...
};

struct FunctionSymbol {
    size_t labelPos{0};
    size_t arity{0};
    std::vector<core::Primitive> paramTypes;
    core::Primitive returnType;

    ...
};
```

Listing 32: Definitions of `VariableSymbol` and `FunctionSymbol` (backend/Symbol.hpp)

Again the additional fields present, see Listing 32, within said symbols helps reduce complexity further down the pipeline especially in code generation.

Additionally, it is important to note that only primitives are comparable. Of course being able to compare the full symbols or adding said logic in the `Symbol` class goes against separation of behaviour and state.

Symbol Registration & Search

A Symbol in `PArL` is an attribute associated with an identifier (a string). One of the main jobs of semantic analysis is to solidify these associations from the AST into the current Environment. In `PArL` the only nodes which are meant to cause a registration are the Variable Declaration node and the Function Declaration node.

The only important considerations for registration are that a identifier is never registered twice with in the same scope and that function identifiers are globally unique, that is they cannot be shadowed.

```cpp
Environment *enclosingEnv = env->getEnclosing();

for (;;) {
    if (enclosingEnv == nullptr)
        break;

    std::optional<Symbol> identifierSignature =
        enclosingEnv->findSymbol(param->identifier);

    if (identifierSignature.has_value() &&
        identifierSignature->is<FunctionSymbol>()) {
        error(
            param->position,
            "redeclaration of {}(...) as a "
            "parameter",
            param->identifier
        );
    }
}
```

41

```
        enclosingEnv = enclosingEnv->getEnclosing();
    }

    env->addSymbol(param->identifier, {type});
}

void AnalysisVisitor::visit(core::FunctionDecl *stmt) {
    bool isGlobalEnv = mEnvStack.isCurrentEnvGlobal();

    if (!isGlobalEnv) {
        error(
            stmt->position,
            "function declaration {}(...) is not "
            "allowed "
            "here",
            stmt->identifier
        );
    }
```

Listing 33: The `visit(FormalParam *)` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

Searching happens when a identifier in a none-declaration node is met. The procedure for searching basically requires querying the current environment for the symbol associated with the found identifier. If nothing is found climb to the enclosing state and repeat the process. If the root node or *terminating node* is reached and no associated symbol is found that means that the identifier is undefined.

Note that the specification of a `terminating node` (a `Environment*`) is important. This is because it allows the analyser to restrict access to outer scopes when necessary. The main example of this is function declarations. In PArL functions are (for the most) part pure, that is they do not produce side-effect at a language level. To enforce this, functions are not allowed to access any state outside of their own scope. The only exception to this fact is the usage of built-ins like `__-write` which directly effect the simulator.

This is where one of the extra fields mentioned in 3.2, specifically `mType` is used because it allows scopes to be of different types and this is necessary to find function scopes. For example in Listing 34 `findEnclosingEnv()` is being used to find the closest enclosing function scope and if no such scope is found i.e. the optional is empty. Then searching either terminates on the stopping scope which

42

is either global or the enclosing function scope.

```cpp
void AnalysisVisitor::visit(core::Variable *expr) {
    std::optional<Environment *> stoppingEnv =
        findEnclosingEnv(Environment::Type::FUNCTION);

    std::optional<Symbol> symbol{findSymbol(
        expr->identifier,
        stoppingEnv.has_value() ? *stoppingEnv
                                : mEnvStack.getGlobal()
    )};

    if (!symbol.has_value()) {
        error(
            expr->position,
            "{} is undefined",
            expr->identifier
        );
    }

    ...
```

Listing 34: The `visit(Variable *)` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

3.5 │ Type Checking

This is the more tedious aspect of semantic. Each individual case where the types affect the behaviour of program or are simply not allowed have to be considered. For this the `AnalysisVisitor` has the field `mReturn` which is essentially used as the return of a visit, see Listing 35.

```cpp
void AnalysisVisitor::visit(core::Binary *expr) {
    expr->left->accept(this);
    auto leftType{mReturn};

    expr->right->accept(this);
    auto rightType{mReturn};

    ...
```

43

Listing 35: A part of the `visit(Binary *)` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

In terms of design choice a strict approach was chosen, that is no implicit casting takes place. Operations are not capable of operating with two distinct types, for example `5 / 2.0` is not allowed. Instead operations which should support multiple types have been overloaded (later on in code generation). For example, `__print 5 / 2` outputs 2, `__print 5 / 2.0` is a type error and `__print 5.0 / 2.0` outputs 2.5.

```cpp
...

        case core::Operation::ADD:
        case core::Operation::SUB:
            if (leftType.is<core::Array>()) {
                error(
                    expr->position,
                    "operator {} is not defined on array "
                    "types",
                    core::operationToString(expr->op)
                );
            }

            if (leftType != rightType) {
                error(
                    expr->position,
                    "operator {} expects both operands to "
                    "be of same type",
                    core::operationToString(expr->op)
                );
            }

            mReturn = leftType;
            break;

        ...
```

Listing 36: Another part of the `visit(Binary *)` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

Of course the remainder of the `AnalysisVisitor` is essentially, type checking in the form of described in Listing 36.

Return Validation

```
...

std::optional<Environment *> optEnv =
    findEnclosingEnv(Environment::Type::FUNCTION);

if (!optEnv.has_value()) {
    error(
        stmt->position,
        "return statement must be within a "
        "function block"
    );
}

...
```

Listing 37: Checking whether return statement is inside a function declaration in the `visit(ReturnStmt *)` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

The only non-trivial check is that of a return statement. First of all return statements must be enclosed within functions, otherwise it is a semantic error, see Listing 37.

```
...

Environment *env = *optEnv;

std::string enclosingFunc = env->getName().value();

auto funcSymbol = env->getEnclosing()
                    ->findSymbol(enclosingFunc)
                    ->as<FunctionSymbol>();

if (exprType != funcSymbol.returnType) {
    error(
        stmt->position,
```

```
        "incorrect return type in function {}",
        enclosingFunc
    );
}

...
```

Listing 38: Checking whether the return expression has the same type as the function return type in the `visit(ReturnStmt *)` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

Additionally, they must have the same return type as the enclosing function, see Listing 38, and finally and most importantly, **a function must return in all possible branches**. Additionally, the implementation **should be capable of recognising unreachable areas of code**.

> ❗ Not stopping when an unreachable code segment is met could result in an invalid assessment of whether a function returns from all branches or not.

During the first iteration of the compiler, the mechanism for checking the return type was embedded directly into the semantic analysis. However, their is slight problem with this approach. If an unreachable segment of code is met to avoid the issue mentioned above the remainder of the unreachable code will not be type checked. Of course this is a problem since such behaviour means the compiler can ignore semantically incorrect code and still produce VM instructions.

To solve this issue a separate visitor `ReturnVisitor` was created to handle checking that functions returns from all possible branches. The separation facilitates checking after type checking is complete circumventing any of the discussed issues.

The mechanism which `ReturnVisitor` uses to establish whether a function returns or not is using a class member `mBranchReturns`.

This member is a boolean and it is explicitly set to true only by a return statement. Function declarations reset the `mBranchReturns` to false and other statements which do not exhibit branching do not modify the value of `mBranchReturns`. This means that the only statements which effect `mBranchRetruns` are if-else statements, for-loops and while-loops, since they branch.

For-loops and while-loops reset `mBranchReturns` to false after their block of statements. This is because entry into the body of the loop (both while and for) is conditional and hence a program might never enter into the body of the loop. Hence,

it is not enough to return from the body of the loop.

```cpp
void ReturnVisitor::visit(core::IfStmt *stmt) {
    stmt->thenBlock->accept(this);
    bool thenBranch = mBranchReturns;

    bool elseBranch = false;

    if (stmt->elseBlock) {
        stmt->elseBlock->accept(this);
        elseBranch = mBranchReturns;
    }

    mBranchReturns = thenBranch && elseBranch;
}
```

Listing 39: The `visit(IfStmt *)` method in the `ReturnVisitor` class (analysis/ReturnVisitor.cpp)

If-else statements are different in that if the 'then' block returns and the 'else' block returns then the whole if-else statement returns otherwise the if-else statement does not return. This is essentially a logical and between the results of the visitor for the 'then' and 'else' blocks, see Listing 39.

```cpp
void ReturnVisitor::visit(core::Block *block) {
    size_t i = 0;

    for (; i < block->stmts.size(); i++) {
        block->stmts[i]->accept(this);

        if (mBranchReturns) {
            i++;

            break;
        }
    }

    if (i < block->stmts.size()) {
        core::Position start = block->stmts[i]->position;

        warning(
            start,
```

```
        "statements starting from line {} upto line {} "
        "are unreachable",
        start.row(),
        block->position.row()
    );
  }

  for (; i < block->stmts.size(); i++) { block->stmts.pop_back(); }
}
```

Listing 40: The `visit(Block *)` method in the `ReturnVisitor` class (analysis/ReturnVisitor.cpp)

Finally, in the block statements if one of the statements inside manages to successfully return, the remaining statements will never be reached and hence the user can be notified and the statements can be dropped, see Listing 40.

The main

To do this

This minor change removes the need to throw away Environments and they can be passed on from one phase to another.
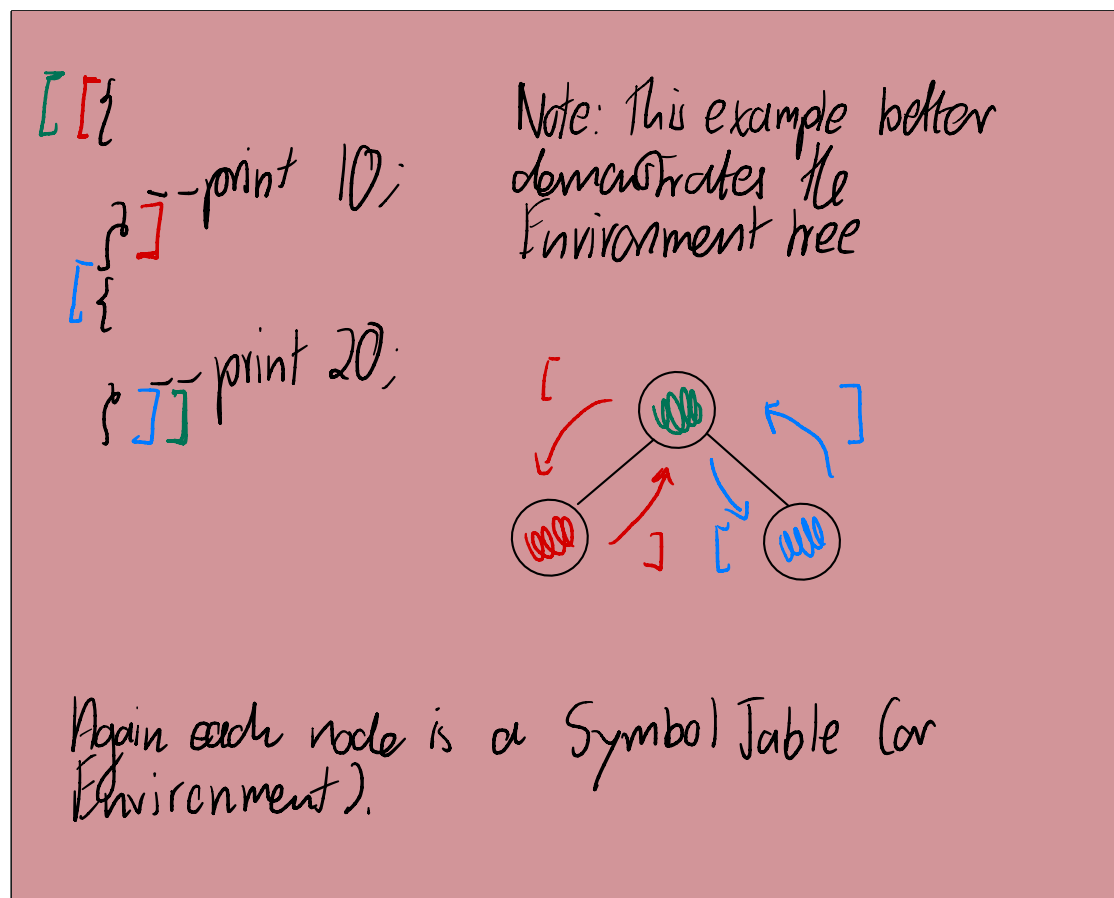
Figure 18: Construction of an `Environment` tree

The Environment tree is a data structure

- Describe Symbol Table.

- Describe each element of the symbol table.

- Describe the tree approach.

- attribute the tree approach to Robert Nystrom as well.

- describe the main advantage/disadvantage of using a environment approach.

- describe the main advantage/disadvantage of using a symbol stack approach.

4 | Code Generation

Having completed semantic analysis, the compiler can proceed to code generation.

4.1 | Passing on the Environment Tree

Before, proceeding onto to the details of code generation it is important to note that here is were all the upfront work that was put into making Environment trees, the data structure of choice for managing scopes, pays off. This is because there is no need to re-stablish the types of any of the variables, at any scope in a program.

```
std::unique_ptr<Environment> environment =
    mAnalyser.getEnvironment();

ReorderVisitor reorder{};

reorder.reorderAst(ast.get());

if (mParserDbg) {
    debugParsing(ast.get());
}

reorder.reorderEnvironment(environment.get());

GenVisitor gen{environment.get()};
```

Listing 41: Getting the global environment from the `AnalysisVisitor` and passing it on for reordering and code generation (runner/Runner.cpp)

The `TypeVisitor` Class

Having said that there is still the need to recompute the types of expressions since the type of a composite expression e.g. `2 + 3` cannot be stored in an Environment. To handle this another visitor, this time called `TypeVisitor` was created.

> **TODO** In the event that type checking/generation is completely delegated to the `TypeVisitor` future version of compiler can drop usage of Environment trees in favour of re-computing types as requested. Additionally, such a change improves memory usage in exchange for time as described in 3.2.

```
void TypeVisitor::visit(core::FunctionCall *expr) {
    std::optional<Symbol> symbol{
        findSymbol(expr->identifier, mRefStack.getGlobal())
    };

    mReturn = symbol->as<FunctionSymbol>().returnType;

    expr->core::Expr::accept(this);
}
```

Listing 42: The `visit(FunctionCall *)` method in the `TypeVisitor` class (ir_gen/TypeVisitor.cpp)

Currently, a `TypeVisitor` is given a `RefStack` for environment access and its main capability is re-calculating the type of **already** type-checked expressions, see Listing 42.

> **!** Usage of a `TypeVisitor` on a non-type-checked expression will at best crash and at worst, still work, leading to undefined behaviour.

4.2 │ Reordering

Due to the fact that the `PArDis` VM Simulator expects all functions to be defined before the `.main` segment, which in the case of `PArL` is everything not in a function declaration, a solution needs to be devised to ensure that this order requirement is satisfied.

A solution to this problem can be achieved by reordering the AST and Environments. This is because function declarations can only exist in global scope, and hence moving them such that they are the first statements, which are encountered in a program is a sufficient constraint to ensure code generation as required by the VM.

> **TODO** Currently the compiler has support for named scopes (see backend/Environemnt.cpp, lines 25 and 34). With a specially designed naming scheme to uniquely identify each scope, support for nested functions can be enable by mangling function names with scope names and lifting function declarations into global scope.

This approach required two visitors, named `IsFunctionVisitor` and `ReorderVisitor`, and a function `reorderEnvironment()` (the choice to embed it into the `ReorderVisitor`

was arbitrary).

The `IsFunctionVisitor` is quite self-explanatory it is a very simple visitor with an exposed method `check()` which returns true only if a node is a function declaration.

```
void ReorderVisitor::reorder(
    std::vector<std::unique_ptr<core::Stmt>> &stmts
) {
    for (auto &stmt : stmts) {
        if (isFunction.check(stmt.get())) {
            mFuncQueue.push_back(std::move(stmt));
        } else {
            mStmtQueue.push_back(std::move(stmt));
        }
    }

    stmts.clear();

    for (auto &stmt : mFuncQueue) {
        stmts.push_back(std::move(stmt));
    }

    for (auto &stmt : mStmtQueue) {
        stmts.push_back(std::move(stmt));
    }

    reset();
}
```

Listing 43:   The `reorder()` method in the `ReorderVisitor` class (preprocess/ReorderVisitor.cpp)

The `ReorderVisitor` implements a single method `reorder()` which is called only for program and block nodes. The `reorder()` makes use of two queues. The first queue keeps track of function declarations and the second queue keeps track of everything else. The `reorder()` method traverse the provided statement vector and enqueues statements in there respective queues. Then the elements of the first queue are dequeued back into the vector and then the elements of the second queue are dequeued back into the vector. Due to the usage of queues this essentially moves all function declarations in their original order as the first elements of the statements vector. The visitor is then accepted by all the reordered

statements.

The visitor is accepted by the reordered statements, to cater for the case when function declarations are also present in other scopes not just the global scope. However, this is currently not the case (see analysis/AnalysisVisitor.cpp, line 955).

4.3 | Function Declarations

Due to the reordering described in 4.2, the first nodes to be emitted are function declarations.

```cpp
void GenVisitor::visit(core::FunctionDecl *stmt) {
    Environment *nextEnv = mRefStack.peekNextEnv();

    size_t aritySize{0};

    for (auto &stmt : stmt->params) {
        aritySize +=
            mDeclCounter.count(stmt.get(), nextEnv);
    }

    emit_line(".{}", stmt->identifier);

    mRefStack.pushEnv(aritySize);

    for (auto &param : stmt->params) {
        param->accept(this);
    }

    stmt->block->accept(this);

    mRefStack.popEnv();
}
```

Listing 44: The `visit(FunctionDecl *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

Again for code generation two more visitors apart from the `TypeVisitor` were used. These are the `GenVisitor` and the `VarDeclCountVisitor`.

```cpp
void VarDeclCountVisitor::visit(core::VariableDecl *stmt) {
    std::optional<Symbol> symbol =
        mEnv->findSymbol(stmt->identifier);

    auto &variable = symbol->asRef<VariableSymbol>();

    mCount += variable.type.is<core::Array>()
                ? variable.type.as<core::Array>().size
                : 1;
}
```

Listing 45: The `visit(VariableDecl *)` method in the `VarDeclCountVisitor` class (ir_gen/VarDeclCountVisitor.cpp)

The `GenVisitor` as its name suggest is the main visitor which produces the VM instructions. The `VarDeclCountVisitor` is an additional visitor which is used to calculate the size of the memory required when opening a frame on the VM.

> **NOTE** Within the VM memory is implemented in the form a JavaScript array. This means that there are no restrictions on how big a single location is. For this reason there is not true standard unit of size. However, it is implicitly assumed that the base types: `bool`, `int`, `float` and `color` all occupy single unit. Hence, when calculating the size of for example `int[4]` the `VarDeclCountVisitor` returns 4 (see Listing 45).

```cpp
template <typename... T>
void emit_line(fmt::format_string<T...> fmt, T &&...args) {
    mCode.push_back(fmt::format(fmt, args...));
}
```

Listing 46: The `emit_line()` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

On the other hand, the main mechanism with which the `GenVisitor` emits VM code is `emit_line()`. Again similar to `error()` and `warning()` from other visitors, `emit_line()` wraps an `fmtlib` function to allow for a nicer overall developer experience when formatting strings, see Listing 44.

In the context of function declarations the only significant usage of `emit_line()` is to output the identifier of the function.

4.4 | Scopes, Variable (or Formal Parameter) Declarations & Variable Accesses

```cpp
void GenVisitor::visit(core::VariableDecl *stmt) {
    stmt->expr->accept(this);

    Environment *env = mRefStack.currentEnv();

    Environment *stoppingEnv = env;

    while (stoppingEnv->getType() !=
                Environment::Type::GLOBAL &&
            stoppingEnv->getType() !=
                Environment::Type::FUNCTION) {
        stoppingEnv = stoppingEnv->getEnclosing();
    }

    std::optional<Symbol> left{};

    for (;;) {
        left = env->findSymbol(stmt->identifier);

        if (left.has_value() || env == stoppingEnv)
            break;

        env = env->getEnclosing();
    }

    core::abort_if(
        !left.has_value(),
        "symbol is undefined"
    );

    size_t idx = env->getIdx();

    auto symbol = left->as<VariableSymbol>();

    if (symbol.type.is<core::Base>()) {
        env->incIdx();
    } else if (symbol.type.is<core::Array>()) {
        env->incIdx(symbol.type.as<core::Array>().size);
    } else {
```

```
            core::abort("unknown type");
    }

    // NOTE: make sure this is actually a reference.
    env->getSymbolAsRef(stmt->identifier)
        .asRef<VariableSymbol>()
        .idx = idx;

    if (symbol.type.is<core::Array>()) {
        emit_line(
            "push {}",
            symbol.type.as<core::Array>().size
        );
    }
    emit_line("push {}", idx);
    emit_line("push 0");
    if (symbol.type.is<core::Array>()) {
        emit_line("sta");
    } else {
        emit_line("st");
    }
```

Listing 47: The `visit(VariableDecl *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

Next, specific attention has to be given to variable declarations and formal parameters. Referencing of formal parameters or variables within the VM is done with the instruction `push [i:l]` (for behaviour see Figure 19). The key takeaway from how this instruction works is the fact that frames (which are actually just memory) are accessed via an index (in the instruction `i`).

Hence, during code generation any reference to a variable or formal parameter needs to happen through an index. This is why, when defining a new variable or formal parameter, the internal index of the current environment is copied into the variable or formal parameter's associated symbol. This facilitates referencing a variable or formal parameter in the VM code (see Listing 47 and ir_gen/GenVisitor.cpp, line 105).

Additionally, the index of the environment is incremented to the next free position i.e. the next position a new variable or formal parameter can be stored within the frame.

Figure 19: A pictorial example of how the `push [i:l]` operates

4.5 | Properly Closing Scopes in Function Declarations

A very important consideration when it comes to function declarations is closing of scopes.

This is a problem of interest because functions should be capable of returning anywhere within the body of a function even if they are nested within multiple scopes.

```
...

emit_line("push {}", count);
emit_line("oframe");

mFrameDepth++;

...
```

Listing 48: A segment of the `visit(Program *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

```cpp
void GenVisitor::visit(core::ReturnStmt *stmt) {
    stmt->expr->accept(this);

    for (size_t i = 0; i < mFrameDepth; i++) {
        emit_line("cframe");
    }

    emit_line("ret");
}
```

Listing 49: The `visit(ReturnStmt *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

To solve this issue, a variable `mFrameDepth` keeps track of the number of frames open with the `oframe` instruction, see Listing 48. When a return statement node is reached the close frame instruction, `cframe`, is emitted as many times as `mFrameDepth`, see Listing 49. This ensures that extra open frames are closed before returning, and hence the VM can continue executing the program normally.

4.6 │ Assignments

Similar, to when variables are first declared the store (`st`) instruction is used for mutating the value in the specified location. However, unlike variable declarations which always define a level of 0, assignments/mutations can occur on variables defined in outer scopes. Hence, the level of variable access needs to be calculated before being emitting `st` and it necessary operands.

```cpp
size_t GenVisitor::computeLevel(Environment *stoppingEnv) {
    size_t level = 0;

    Environment *env = mRefStack.currentEnv();

    while (stoppingEnv != env) {
        switch (env->getType()) {
            case Environment::Type::FOR:
            case Environment::Type::BLOCK:
                level++;
                break;
            case Environment::Type::FUNCTION:
                core::abort("unreachable");
                break;
```

```
            default:
                // noop
                break;
        }

        env = env->getEnclosing();
    }

    return level;
}
```

Listing 50: The `computeLevel()` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

The basic idea for calculating the level of a variable is as follows, copy the pointer to the environment which holds the variable then traverse from the current environment until the environment with the exact pointer as the environment the variable was found in is reached, all the while keeping track of the number of steps scopes jumped through to reach said environment.

> **NOTE**
> Not, all environments contribute an open frame (`oframe`), hence only certain types of scopes should be counted as contributing to the number of levels. In the case of `PArL` with its current restrictions, the only such scopes are for-scopes and block-scopes (see Listing 50).

4.7 │ Expressions

Type Specific VM Code

```
...
case core::Operation::DIV: {
    core::Primitive type = mType.getType(
        expr->right.get(),
        mRefStack.getGlobal(),
        mRefStack.currentEnv()
    );
    if (type == core::Primitive{core::Base::INT}) {
        expr->right->accept(this);
        expr->right->accept(this);
        expr->left->accept(this);
        emit_line("mod");
```

```
        expr->left->accept(this);
        emit_line("sub");
        emit_line("div");
    } else {
        expr->right->accept(this);
        expr->left->accept(this);
        emit_line("div");
    }
} break;
...
```

Listing 51: A segment of the `visit(Binary *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

When it comes to expressions, due to the fact that operators within the PArL language are overloaded some attention to detail is required. The majority of the usage of `TypeVisitor` is within expression nodes, specifically, in places where behaviour is dependent on operand types. For example in the case of division, the type of one of the operands is computed (recall that the semantic analysis phase guarantees consistency of types, in the case of binary operations both the left and right operands have the same type), then different instructions are emitted depending on the returned type, see Listing 51.

Function Calls

```
void GenVisitor::visit(core::FunctionCall *expr) {
    for (auto itr = expr->params.rbegin();
         itr != expr->params.rend();
         itr++) {
        (*itr)->accept(this);
    }

    size_t size{0};

    for (auto &param : expr->params) {
        core::Primitive paramType = mType.getType(
            param.get(),
            mRefStack.getGlobal(),
            mRefStack.currentEnv()
        );
```

60

```
        size += paramType.is<core::Array>()
                    ? paramType.as<core::Array>().size
                    : 1;
    }

    emit_line("push {}", size);
    emit_line("push .{}", expr->identifier);
    emit_line("call");
}
```

Listing 52: The `visit(FunctionCall *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

Functions calls are performed using the `call` VM instruction. The most important aspect of function calls, similar to scopes, is the calculation of the total size of the parameters. This essentially determines how many operand will pop of the operand stack and moved into the implicitly created function frame. Additionally, it is important for all the parameters of a function to be evaluated in reverse order due to the fact that popping from the operand stand into the function frame essentially reverse the order of the parameters (see Listing 52).

4.8 │ Branching

Branching is somewhat challenging since at the jumping point the compiler is unaware of how many instructions it might need to jump. There are two possible solutions for this problem. The first is to create a visitor which is capable of calculating the number of instruction which a node expands to. The second makes use of a placeholder instruction that is then changed after the rest of the associated nodes are expanded.

Previous iterations of the `PArL` compiler used the first approach however, the first approach is quite unmaintainable since changes in other areas of the code generation would require changes in the visitor which computes the number of instructions. Hence, the compiler was shifted to the second approach.

`PArL` has only three node types which branch, if-else statements, for-loops and while-loops.

```
void GenVisitor::visit(core::WhileStmt *stmt) {
    mRefStack.pushEnv();
```

```cpp
    size_t condOffset = PC();

    stmt->cond->accept(this);

    emit_line("not");

    size_t patchOffset = PC();

    emit_line("push #PC+{{}}");
    emit_line("cjmp");

    stmt->block->accept(this);

    emit_line("push #PC-{}", PC() - condOffset);
    emit_line("jmp");

    mCode[patchOffset] =
        fmt::format(mCode[patchOffset], PC() - patchOffset);

    mRefStack.popEnv();
}
```

Listing 53: The `visit(WhileStmt *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

Both for-loops and while-loops branch in the same manner. Patching the relevant instructions for branching in for- and while-loops works as follows (additionally, see Listing 53):

1. The location of the instruction which starts the condition is stored, in this case in `condOffset`

2. the location of the `push #PC+{}` instruction used for `cjmp` is stored, in this case in `patchOffset`;

3. then the block associated with the for- or while-loop accepts the `GenVisitor`, emitting further instructions

4. then the offset required to jump back to the start of the condition instructions is calculated and used with a `jmp` instruction;

5. finally, the `push #PC+{}` instruction at `patchOffset` is patched with the appropriate offset to jump clear all the for or while block and the instructions for looping.

```
void GenVisitor::visit(core::IfStmt *stmt) {
    if (stmt->elseBlock) {
        mRefStack.pushEnv();

        stmt->cond->accept(this);

        emit_line("not");

        size_t ifPatchOffset = PC();

        emit_line("push #PC+{{}}");
        emit_line("cjmp");

        stmt->thenBlock->accept(this);

        mRefStack.popEnv();

        mRefStack.pushEnv();

        size_t elsePatchOffset = PC();

        emit_line("push #PC+{{}}");
        emit_line("jmp");

        mCode[ifPatchOffset] = fmt::format(
            mCode[ifPatchOffset],
            PC() - ifPatchOffset
        );

        stmt->elseBlock->accept(this);

        mCode[elsePatchOffset] = fmt::format(
            mCode[elsePatchOffset],
            PC() - elsePatchOffset
        );

        mRefStack.popEnv();
    } else {
```

```
        mRefStack.pushEnv();

        stmt->cond->accept(this);

        emit_line("not");

        size_t patchOffset = PC();

        emit_line("push #PC+{{}}");
        emit_line("cjmp");

        stmt->thenBlock->accept(this);

        mCode[patchOffset] = fmt::format(
            mCode[patchOffset],
            PC() - patchOffset
        );

        mRefStack.popEnv();
    }
}
```

Listing 54: The `visit(IfStmt *)` method in the `GenVisitor` class (ir_gen/GenVisitor.cpp)

The other form of branching, if-else statements, is a bit more involved. This is because there are two cases which need consideration: when only an `if` is present and when both an `if` and an `else` are present. Nevertheless, the basic principal is the same offsets into the code are used and specific lines are patched after then fact (see Listing 54).

5 | Arrays

Adding support for arrays in `PArL` has driven a number of changes across whole of the compiler. Some of these changes will be explained in the following sections whilst otherwise where already described in earlier section due to the fundamental changes they caused.

5.1 | Changes in Lexical Analysis

```
...

.addCategory(
    LEFT_BRACKET,
    [](char c) -> bool {
        return c == '[';
    }
)
.addCategory(
    RIGHT_BRACKET,
    [](char c) -> bool {
        return c == ']';
    }
)

...

.addTransition(0, LEFT_BRACKET, 20)
.addTransition(0, RIGHT_BRACKET, 21)

...

.setStateAsFinal(20, Token::Type::LEFT_BRACK)
.setStateAsFinal(21, Token::Type::RIGHT_BRACK)

...
```

Listing 55: Changes in the `LexerDirector` for supporting left ('[') and right (']') brackets (lexer/LexerDirector.cpp)

Within lexical analysis a very lax approach to arrays was adopted. That is there is no enforcement of the more complicated ⟨Identifier⟩ syntax proposed in the original EBNF of the language. The only additions have been support for consuming left ('[') and right (']') brackets.

5.2 | Changes in Parsing

As described in 2.1, the grammar of the language was significantly reordered to improve the upgrade-ability of the language down the line. Additionally, the majority of the changes to the parser where also discussed in 2.1.

> **TODO**
> The addition of de-sugaring, type inference and more types is facilitated by the changes proposed in the EBNF. These additions would allow for more complicated language mechanisms and hence an improved developer experience down the line.

5.3 | Changes in Semantic Analysis

Again due to the changes in the way types are managed, (again see 2.1), this leads to a number of notable changes in particular: extra validation is necessary, restriction on type casting, figuring out the type of array access and finally, generating the type for array literals.

```
...
case core::Operation::ADD:
case core::Operation::SUB:
    if (leftType.is<core::Array>()) {
        error(
            expr->position,
            "operator {} is not defined on array "
            "types",
            core::operationToString(expr->op)
        );
    }

    if (leftType != rightType) {
        error(
            expr->position,
            "operator {} expects both operands to "
            "be of same type",
            core::operationToString(expr->op)
        );
    }

    mReturn = leftType;
    break;
...
```

Listing 56: A part of the `visit(Binary *)` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

Of course validation is necessary to ensure that arrays do no interact with each

other, for example `[1,2,3] + [1,2,3]` is not semantically, correct in `PArL`.

```cpp
void AnalysisVisitor::isViableCast(
    core::Primitive &from,
    core::Primitive &to

) {
    const core::Primitive *lPtr = &from;
    const core::Primitive *rPtr = &to;

    while (lPtr != nullptr && rPtr != nullptr) {
        if (lPtr->is<core::Base>() &&
            rPtr->is<core::Array>()) {
            error(
                mPosition,
                "primitive cannot be cast to an array"
            );
        }

        if (lPtr->is<core::Array>() &&
            rPtr->is<core::Base>()) {
            error(
                mPosition,
                "array cannot be cast to a primitive"
            );
        }

        if (lPtr->is<core::Array>() &&
            rPtr->is<core::Array>()) {
            size_t lSize = lPtr->as<core::Array>().size;
            size_t rSize = rPtr->as<core::Array>().size;

            lPtr = lPtr->as<core::Array>().type.get();
            rPtr = rPtr->as<core::Array>().type.get();

            if (lSize != rSize) {
                error(
                    mPosition,
                    "array of size {} cannot be cast to an "
                    "array of size {}",
                    lSize,
                    rSize
                );
```

```
        }

        continue;
    }

    return;
    }

    core::abort("unreachable");
}
```

Listing 57: The `isViableCast()` method in the `AnalysisVisitor` class (analysis/AnalysisVisitor.cpp)

Necessary checks to ensure that casting is viable, need to be considered. This approach was taken over allowing all forms of casting because casting from one type to another where the sizes of the types are different would result in substantial complexity later on in code generation.

Consider the following example:

```
let x: int[5] = [1,2,3,4,5];
let y: int = (x as color[10])[6];
```

Not only is there ambiguity as to what the upper elements of the type cast should be but the additional elements are necessary to ensure no out of of bounds accesses during VM execution.

Finally, the process of computing the type of an array access and an array literal is quite easy, see analysis/AnalysisVisitor.cpp lines 252 and 170 respectively.

5.4 | Changes in Code Generation

One of the more important changes for code generation is the calculation of frame sizes. This has already been described in detail even in the context of arrays in the note in subsection **??**. Of course where appropriate `sta`, `pusha`, `printa` and `push +[i:l]` which are for storing arrays, pushing arrays onto the operand stack, printing arrays and pushing an single element of the array onto the operand

stack respectively, are being used accordingly. See ir_gen/GenVisitor.cpp, lines 456, 111, 478 and 166, for examples of each.

A decision to opt out of using `reta` was made due to the inconsistent behaviour of arrays in the VM. Instead a hack for reversing the order manually was devised.

```cpp
...
if (symbol.type.is<core::Array>()) {
    size_t arraySize =
        symbol.type.as<core::Array>().size;
    emit_line("push {}", arraySize);
    emit_line("pusha [{}:{}]", symbol.idx, level);
    emit_line("push {} // START HACK", arraySize);
    emit_line("oframe");
    emit_line("push {}", arraySize);
    emit_line("push 0");
    emit_line("push 0");
    emit_line("sta");
    emit_line("push {}", arraySize);
    emit_line("pusha [0:0]");
    emit_line("cframe // END HACK");

    return;
}
...
```

Listing 58: A part of the `visit(Variable *)` method in the `GenVisitor` class, specifically the hack mentioned above (ir_gen/GenVisitor.cpp)

The problem of arrays is the following:

Consider the following code: `__print [1,2,3];`

```
.main
push 0
oframe
push 3
push 2
push 1
push 3
printa
cframe
halt
```

It generates the code on the left and what is important to notice is the last element in the array is the first to be pushed onto the operand stack ensuring that printing is correct.

However, now consider the following piece of code: `let x: int[3] = [1,2,3];` `__print x;`

```
.main
push 3
oframe
push 3
push 2
push 1
push 3
push 0
push 0
sta
push 3
pusha
  [0:0]
push 3
printa
cframe
halt
```

The nïve implementation generates the code on the left. Running this on the VM one would expect the exact same output as the first. However, this is not the case the order of the elements is reversed.

```
.main
push 3
oframe
push 3
push 2
push 1
push 3
push 0
push 0
sta
push 3
pusha [0:0]
push 3
printa
cframe
halt
```

=>>> logs start
here
-- 3
-- 2
-- 1

Figure 20: The result of running the code generated `let x: int[3] = [1,2,3];
__print x;` in the VM

This happens because reading from a frame using `pusha` actually reverse the order of the elements on the operand stack. To circumvent this issue a second temporary frame is created where the elements are stored in the frame and read again reversing the order one-more time into the correct order (see Listing 58).

Due to this change hack which works in all situations `reta` is not being used as the would flip the order again.

6 | Testing & Usage

6.1 | Usage

The current produced binary accepts the following flags `-d`, `-l`, `-p`. `-d` is currently disabled, `-l` prints each of the lexemes generated by the source and `-p` outputs the AST both before and after preprocessing.
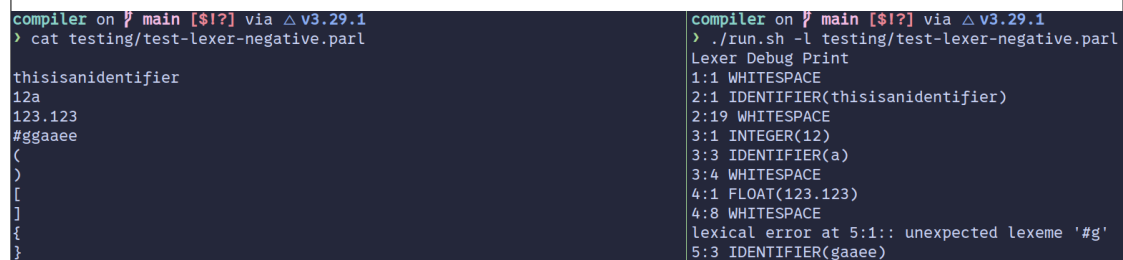
6.2 | The Lexer, Parser & Semantic Analyser

As discussed in previous sections all the phases which are capable of producing errors or warnings are either doing so manually using `fmt::println(stderr,`

`...)` or using some customised methods like `error(...)` or `warning(...)`. Additionally, as described in previous sections as well, the lexer, parser and analyser should not fail upon the first inconsistency within the program, instead they should proceed to a state in which they continue processing normally.

The main aim of testing will be to make sure that both correct and incorrect input produce the expected behaviour.

Lexical Analysis

```
compiler on  main [$!?] via △ v3.29.1          compiler on  main [$!?] via △ v3.29.1
> cat testing/test-lexer-negative.parl         > ./run.sh -l testing/test-lexer-negative.parl
                                                Lexer Debug Print
thisisanidentifier                             1:1 WHITESPACE
12a                                            2:1 IDENTIFIER(thisisanidentifier)
123.123                                        2:19 WHITESPACE
#ggaaee                                        3:1 INTEGER(12)
(                                              3:3 IDENTIFIER(a)
)                                              3:4 WHITESPACE
[                                              4:1 FLOAT(123.123)
]                                              4:8 WHITESPACE
{                                              lexical error at 5:1:: unexpected lexeme '#g'
}                                              5:3 IDENTIFIER(gaaee)
```
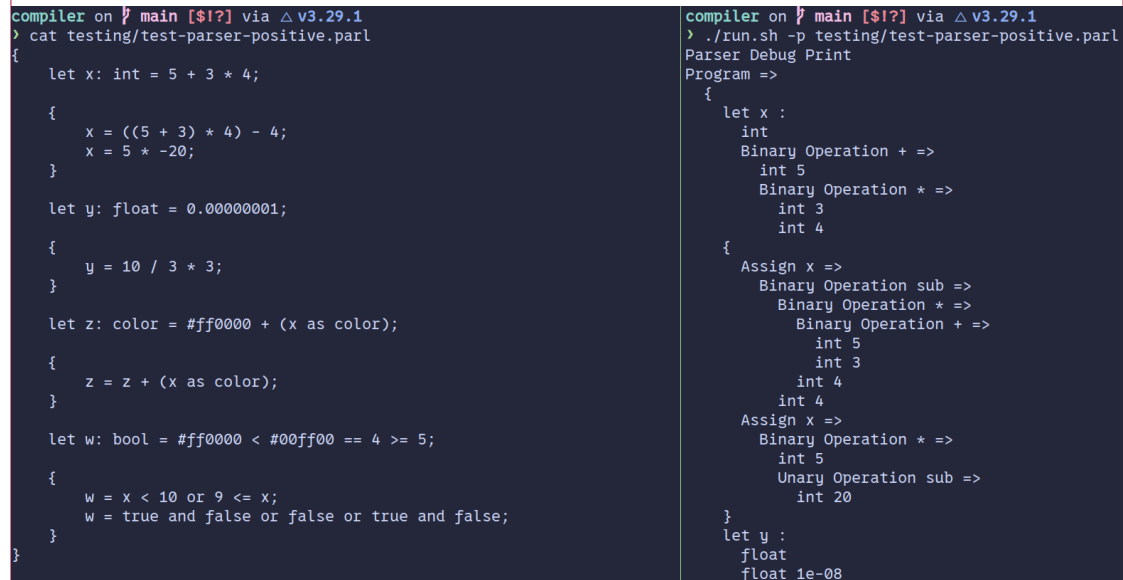
Figure 21: Running the negative lexer test (testing/test-lexer-negative.parl)

For lexical analysis two major tests where used, testing/test-lexer-negative.parl and testing/test-lexer-positive.parl. Additional, tests include, testing/test1.parl upto testing/test9.parl.

Parsing

```
compiler on  main [$!?] via  v3.29.1
> cat testing/test-parser-positive.parl

    let x: int = 5 + 3 * 4;

    {
        x = ((5 + 3) * 4) - 4;
        x = 5 * -20;
    }

    let y: float = 0.00000001;

    {
        y = 10 / 3 * 3;
    }

    let z: color = #ff0000 + (x as color);

    {
        z = z + (x as color);
    }

    let w: bool = #ff0000 < #00ff00 == 4 >= 5;

    {
        w = x < 10 or 9 <= x;
        w = true and false or false or true and false;
    }
}
```
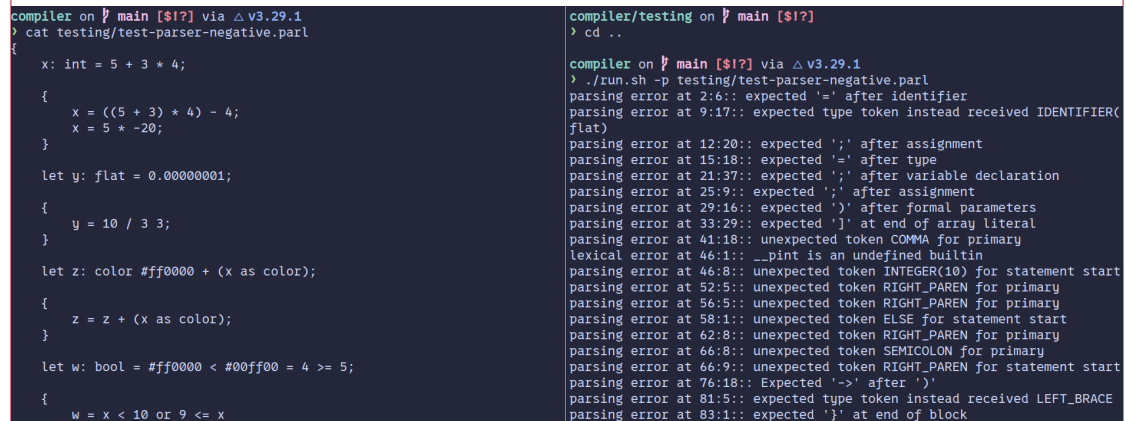
```
compiler on  main [$!?] via  v3.29.1
> ./run.sh -p testing/test-parser-positive.parl
Parser Debug Print
Program =>
  {
    let x :
      int
      Binary Operation + =>
        int 5
        Binary Operation * =>
          int 3
          int 4
      {
        Assign x =>
          Binary Operation sub =>
            Binary Operation * =>
              Binary Operation + =>
                int 5
                int 3
              int 4
            int 4
        Assign x =>
          Binary Operation * =>
            int 5
            Unary Operation sub =>
              int 20
      }
    let y :
      float
      float 1e-08
```

Figure 22: Running the positive parser test (testing/test-parser-positive.parl)

Again for parsing two major tests where developed testing/test-parser-negative.parl and testing/test-parser-positive.parl.

```
compiler on  main [$!?] via  v3.29.1
> cat testing/test-parser-negative.parl
{
    x: int = 5 + 3 * 4;

    {
        x = ((5 + 3) * 4) - 4;
        x = 5 * -20;
    }

    let y: flat = 0.00000001;

    {
        y = 10 / 3 3;
    }

    let z: color #ff0000 + (x as color);

    {
        z = z + (x as color);
    }

    let w: bool = #ff0000 < #00ff00 = 4 >= 5;

    {
        w = x < 10 or 9 <= x
```

```
compiler/testing on  main [$!?]
> cd ..

compiler on  main [$!?] via  v3.29.1
> ./run.sh -p testing/test-parser-negative.parl
parsing error at 2:6:: expected '=' after identifier
parsing error at 9:17:: expected type token instead received IDENTIFIER(
flat)
parsing error at 12:20:: expected ';' after assignment
parsing error at 15:18:: expected '=' after type
parsing error at 21:37:: expected ';' after variable declaration
parsing error at 25:9:: expected ';' after assignment
parsing error at 29:16:: expected ')' after formal parameters
parsing error at 33:29:: expected ']' at end of array literal
parsing error at 41:18:: unexpected token COMMA for primary
lexical error at 46:1:: __pint is an undefined builtin
parsing error at 46:8:: unexpected token INTEGER(10) for statement start
parsing error at 52:5:: unexpected token RIGHT_PAREN for primary
parsing error at 56:5:: unexpected token RIGHT_PAREN for primary
parsing error at 58:1:: unexpected token ELSE for statement start
parsing error at 62:8:: unexpected token RIGHT_PAREN for primary
parsing error at 66:8:: unexpected token SEMICOLON for primary
parsing error at 66:9:: unexpected token RIGHT_PAREN for statement start
parsing error at 76:18:: Expected '->' after ')'
parsing error at 81:5:: expected type token instead received LEFT_BRACE
parsing error at 83:1:: expected '}' at end of block
```

Figure 23: Running the negative parser test (testing/test-parser-negative.parl)

In particular the positive test covers of the happy paths of the parser. However, there are way more possibilities where the parser flags the input as not conforming to the grammar. Hence, not all possible code paths are being tested with

the negative test. Additional, files used for testing during development include: testing/test10.parl, testing/test15.parl and testing/test16.parl.

Semantic Analysis

Semantic analysis is again a leap in complexity and leap in the order of possible errors and issues. Hence, completely testing all possible scenarios is very diffcult. Instead the tested presented and mentioned are those tested written during development to test certain aspects of the semantic analyser as it was being developed.

The tests mentioned above are:

- testing/test13.parl;
- testing/test14.parl;
- testing/test17.parl;
- testing/test18.parl (parameter types);
- testing/test19.parl (parameter types);
- testing/test21.parl (casting);
- testing/test23.parl (no main functions);
- testing/test24.parl (for nested functions);
- testing/test25.parl;
- testing/test28.parl (for unreachable code);
- testing/test31.parl (for arrays);
- testing/test32.parl;
- testing/test33.parl;
- testing/test41.parl (for unreachable code).

7 | Attributions

- Sandro Spina for the brilliant description of table-driven lexers

- Robert Nystrom and his great book Crafting Intepreters for a great outline for parsing and error recovery/management for languages which support exceptions